

# A Tutorial on Gaussian Process Exploration and Exploitation

Eric Schulz\*, Maarten Speekenbrink\*, Andreas Krause\*\*

*Computational Learning and Decision Making Laboratory  
University College London*

*Learning and Adaptive Systems Group  
ETH Zürich*

---

## Abstract

This tutorial introduces Gaussian Process Regression as a method to explore (learn) and exploit (maximize) unknown functions. Gaussian Process Regression is a popular non-parametric Bayesian approach to sequential function learning problems. This tutorial aims to provide an accessible introduction to these techniques. We will introduce and define Gaussian Processes as a distribution over functions used for Bayesian inference and demonstrate different applications thereof. Examples will focus on pure exploration within optimal design problems, bandit-like exploration-exploitation scenarios, and on scenarios with additional constraints, for example safe explorations, where the goal is to never sample below a certain threshold. Software pointers will be provided.

*Keywords:* Gaussian Process, Exploration-Exploitation, Bandit Problems

---

---

\*Department of Experimental Psychology, University College London.

\*\*Department of Computer Science, Swiss Federal Institute of Technology Zürich.

## 1. Introduction

Whether we try to find a function that describes participants' behavior (Cavagnaro, Aranovich, McClure, Pitt, and Myung, 2014), estimate parameters of psychological models (Wetzels, Vandekerckhove, Tuerlinckx, and Wagenmakers, 2010), sequentially optimize stimuli in an experiment (Myung and Pitt, 2009), or model how participants themselves learn to interact with their environment (Meder and Nelson, 2012), many research problems require us to explore (learn) or exploit (optimize) an unknown function that maps inputs to outputs (Mockus, 2010). Often, the underlying function is unknown, hard to evaluate analytically, or other requirements such as design costs might complicate the process of information acquisition. In these situations, Gaussian Process regression can serve as a useful multi-purpose tool towards learning and optimization (Rasmussen, 2006). Gaussian Process regression is a non-parametric Bayesian approach (Gershman and Blei, 2012) to regression problems. It can capture a wide variety of functional input-output relations by utilizing a potentially infinite number of parameters and letting the data “decide” on the level of complexity through Bayesian posterior inference (Williams, 1998).

This tutorial will introduce Gaussian Process Exploration-Exploitation as an approach towards learning and optimizing unknown functions. It consists of 5 main parts: The first part will introduce the mathematical basis of Gaussian process regression. The second part will show how Gaussian processes can be used in problems of optimal experimental design, where the goal – learning a function as well as possible – is purely explorative. The third part will describe Bayesian optimization (*exploitation*) with Gaussian processes. **In the fourth part, we will discuss the use of Gaussian**

27 **process exploration-exploitation in situations with additional re-**  
 28 **quirements and show how they can be applied when the goal is to**  
 29 **avoid areas that are below a certain threshold.** We will conclude by  
 30 sketching out the possibility of Gaussian Processes as a low level description  
 31 of cognitive processes.

## 32 **2. Problem Statement**

33 Imagine a function  $f : x \rightarrow y$  that maps an input  $x$  to an output  
 34  $y$ . Initially, the function is unknown. The task is to choose sequentially,  
 35 at each time or trial  $t$ , an input value as wisely as possible, either to learn  
 36 about the unknown function (exploration) or to find the input for which  
 37 an optimum output is obtained *exploitation*. More formally, at each time  
 38  $t$ , the function is queried by choosing a point  $x_t^*$  for which the output of  
 39 the function  $y_t = f(x_t^*) + \epsilon_t$  at that point is observed with additional noise  
 40  $\epsilon_t$ . It is assumed that this noise is independent and identically distributed,  
 41 following a Gaussian distribution:  $\epsilon_t \sim \mathcal{N}(0, \sigma^2)$ .

### 42 *2.1. Exploration*

43 The goal within exploration is to choose  $x_t^*$  in a way that allows you to  
 44 learn an unknown function as well as possible. This is sometimes referred to  
 45 as an optimal experimental design problem (Goos and Jones, 2011). More  
 46 formally, standing at time  $t$ , the task is to find an informative set of points  
 47  $x_{t:T} = (x_t, \dots, x_T) \subset \prod_{u=t}^T \mathcal{D}_u$ , that is a sequence of queries  $x_t$  chosen  
 48 from the design set  $\mathcal{D}_t$ , chosen to maximize information gain over all future  
 49 times. Generally, the information gain is measured as the expected difference  
 50 in entropy between the state before the sample points were queried and  
 51 afterwards:

$$I(f; x_{t:T}) = E[H(f|x_{1:(t-1)}, y_{1:(t-1)}) - H(f|x_{1:T}, y_{1:(t-1)})] \quad (1)$$

where the expectation is taken over  $p(y_{t:T})$ .

**In the continuous case, the entropy  $H$  can be calculated as**

$$H(f|x_{i:j}) = - \int p(f|x_{i:j}) \log p(f|x_{i:j})$$

53

54 Finding the absolute information gain maximizer is generally NP-hard  
 55 (Ko, Lee, and Queyranne, 1995), which means that one can never know  
 56 the best strategy of how to pick all of the observations over time to cer-  
 57 tainly reach the maximum information gain at the end as this requires an  
 58 exponentially increasing problem space to be evaluated.

59 Optimal design problems are ubiquitous in psychology. No matter if  
 60 we have to find out how participants integrate information (Borji and Itti,  
 61 2013) or if we have to find stimuli in order to distinguish between different  
 62 models (Cavagnaro, Myung, Pitt, and Kujala, 2010), often times we have  
 63 to try and learn an underlying function that we do not know a priori and  
 64 do so both adaptively and efficiently. As we will see later, Gaussian Process  
 65 exploration algorithms provide us with an excellent tool for such tasks.

## 66 *2.2. Exploration-Exploitation*

67 The goal within exploration-exploitation tasks is to both learn and op-  
 68 timize a function. Here, it is essentially the intention to learn a function  
 69 quickly so that we then can find the maximum of that function and keep on  
 70 exploiting it; exploiting means entering the input that produces the highest  
 71 output. This is normally measured by regret, the distance between what

72 your current guess of the maximum has produced and what the best argu-  
 73 ment would have produced. If you are sequentially producing an estimate of  
 74 what you currently think might be the maximum  $x^*$ , then it is possible to  
 75 measure regret  $R$  by the difference to what the actual best argument  $x_{\max}$   
 76 would have produced.

$$R = \sum_{i=1}^T r_i$$

$$= \sum_{i=1}^T f(x_{\max}) - f(x^*)$$

77 Problems where we have to optimize an unknown function are very com-  
 78 mon; be it to estimate parameters of a complex model (Snoek, Larochelle,  
 79 and Adams, 2012) or finding the stimulus that produces the maximal re-  
 80 sponse in an experiment (Daunizeau, Preuschoff, Friston, and Stephan,  
 81 2011), many problems require us to explore and exploit functions. Some-  
 82 times this approach is also called Bayesian Optimization (Brochu, Cora, and  
 83 De Freitas, 2010).

### 84 **3. Gaussian Processes-a distribution over functions**

#### 85 *3.1. Motivation*

86 If our goal is to learn or optimize an unknown function, then we need  
 87 the following two ingredients to be successful.

- 88 1. A model for  $f$ , that is a model that learns what  $f$  looks like.
- 89 2. A method to select observations, given the problem statement.

90 In this section we want to find a good model for  $f$  before we can then  
 91 use it to either explore or exploit the function.

### 92 3.2. Weight space view

93 Let us first start by considering the text book-approach to learn func-  
 94 tions, which is linear Bayesian regression, before then transitioning over to  
 95 Gaussian Process regression. Remember that it is the task to either learn or  
 96 optimize an unknown function. Therefore, regressing observed input points  
 97 to observed output points by standard regression seems to be an apparent  
 98 thing to do. Once we are able to capture the function –so the intuition–  
 99 we can easily quantify our uncertainty about it or try to find the point that  
 100 produces the highest expected output. In classic linear regression, we as-  
 101 sume normally distributed noise,  $\epsilon \sim \mathcal{N}(0, \Sigma)$ , and our goal is to predict a  
 102 value  $y$  based on observations  $\mathbf{x}$  using weights  $\mathbf{w}$ .

103

$$f(\mathbf{x}|\mathbf{w}) = \mathbf{x}^\top \mathbf{w}$$

$$y = f + \epsilon$$

104 Here,  $\mathbf{x}$  are our observations and  $\mathbf{w}$  is a vector containing the weights  
 105 (sometimes notated as  $\beta$ ). If we assume a Gaussian prior over the parameters  
 106  $p(\mathbf{w}) = \mathcal{N}(0, \Sigma)$  and the likelihood  $p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{X}^\top \mathbf{w}, \sigma^2 \mathbf{I})$ , then –by  
 107 applying standard Bayesian inference– we get the posterior

$$\begin{aligned} p(\mathbf{w}|\mathbf{y}, \mathbf{X}) &\propto p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}) \\ &= \mathcal{N}\left(\frac{1}{\sigma^2} \mathbf{A}^{-1} \mathbf{X} \mathbf{y}, \mathbf{A}^{-1}\right) \end{aligned}$$

108 with  $\mathbf{A} = \Sigma^{-1} + \sigma^{-2} \mathbf{X} \mathbf{X}^\top$ . In order to make predictions at a new test  
 109 points  $\mathbf{x}_\star$ , we have to average over all posterior predictions

$$\begin{aligned}
p(f_\star|\mathbf{x}_\star, \mathbf{X}, \mathbf{y}) &= \int p(f_\star|\mathbf{x}_\star, \mathbf{w})p(\mathbf{w}|\mathbf{X}, \mathbf{y})d\mathbf{w} \\
&= \mathcal{N}(\sigma^{-2}\mathbf{x}_\star^\top \mathbf{A}^{-1}\mathbf{X}\mathbf{y}, \mathbf{x}_\star^\top \mathbf{A}^{-1}\mathbf{x}_\star)
\end{aligned}$$

110 Even though this approach is commonly chosen to model functions, the  
111 way it is set up above only allows to make linear predictions. However,  
112 only few relations in the real world are actually linear. Therefore, what  
113 we need is a way to model non-linear dependencies as well. One possible  
114 adjustment is to use a projection of the inputs  $\mathbf{x}$  onto a feature space by using  
115 a function  $\phi(\mathbf{x})$ . One common projection is to use polynomials, resulting  
116 into polynomial regression. Take a cubic regression as an example, which  
117 assumes a function  $f(x) = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \epsilon$ . Deriving the posterior  
118 for this model is similar to the linear regression described before, only that  
119 all  $\mathbf{X}$  are replaced by the projection  $\Phi = \phi(\mathbf{X})$ . However, the problem then  
120 becomes to find appropriate projections for our input variables as infinitely  
121 many projections might be possible and we have to choose one a priori (or  
122 by model comparison between finite sets of parametric forms). Especially if  
123 the problem is to explore and exploit a completely unknown function, this  
124 approach will not be beneficial as we would not know which projections we  
125 should try out (we do not know the parametric form a priori). Fortunately,  
126 there exists another approach to this problem which treats functions as  
127 distributions and models this distribution directly. This approach is called  
128 Gaussian Process regression and shifts the view away from a weight space  
129 perspective to a function space view.

### 130 3.3. Function space view

131 Another approach to regression problems is –instead of modeling distri-  
 132 butions over weights– to focus on distributions over functions. A common  
 133 way to describe a distribution over functions is a Gaussian Process. A  
 134 Gaussian process is defined as a collection of random variables, any finite  
 135 number of which have a jointly Gaussian distribution. Therefore, a function  
 136 is distributed as

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

137 Here,  $m(\mathbf{x})$  is a mean function modeling the expected output of the  
 138 function and  $k(\mathbf{x}, \mathbf{x}')$  is a kernel function modeling the covariance between  
 139 different points.

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

$$k(x, x') = \mathbb{E}[(f(\mathbf{x}) - m(x))(f(\mathbf{x}') - m(x'))]$$

140 As  $k(x, x')$  models the covariance between two different points, we have  
 141 to choose a function that will produce sensible estimates. The function  $k$  is  
 142 commonly called the *kernel* of the Gaussian Process (Jäkel, Schölkopf, and  
 143 Wichmann, 2007). The choice of an appropriate kernel is normally based on  
 144 assumptions such as smoothness and likely patterns to be expected in the  
 145 data. If the data is not expected to be periodic, then a sensible assumption  
 146 is normally that the correlation between two points decays according to a  
 147 power function in dependency of the distance between the two points and  
 148 that the covariance is symmetric, that is that only the distance between two  
 149 points matters, but not the direction. This just means that closer points



are expected to be more similar than points which are further away from each other in all possible directions. One very popular choice of a kernel fulfilling those requirement is the so-called squared exponential (sometimes also called Gaussian or Radial Basis Function) kernel.

$$k(x, x') = \sigma^2 \exp\left(-\frac{(x - x')^2}{2\lambda^2}\right)$$

The squared exponential is an expressive way to model smooth functions and the hyper parameters  $\lambda$  (called the length-scale) and  $\sigma^2$  (the noise constant) are normally optimized by using the marginal likelihood.

This implies the aforementioned distribution over functions as we can easily generate samples for new input points at location  $X_\star$ .

$$\mathbf{f}_\star \sim \mathcal{N}(0, K(X_\star, X_\star))$$

Given observations  $\mathcal{D} = \{\mathbf{X}, \mathbf{y}\}$  with a noise level  $\sigma$ , we can draw new predictions from our function  $\mathbf{f}_\star$  for inputs  $X_\star$  as described below.

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_\star \end{bmatrix} \sim \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} K(X, X) + \sigma^2 I & K(X, X_\star) \\ K(X_\star, X) & K(X_\star, X_\star) \end{bmatrix}\right)$$

This means that we treat a function as a vector of infinite size. However, as we always only have to make predictions for finitely many points, we can simply draw outputs for these points by using a multivariate normal distribution with a covariance matrix generated by our kernel. Calculating the expectation of the Gaussian Process at the new points then is straight forward.

$$\mathbf{f}_\star | X, \mathbf{y}, X_\star \sim \mathcal{N}(\bar{\mathbf{f}}_\star, \text{cov}(\mathbf{f}_\star)) \quad \text{where}$$

167 This means that predictions for new points are generated based on the  
 168 expected mean value and covariance function of the posterior Gaussian Pro-  
 169 cess.

$$\begin{aligned} \mathbb{E}[\mathbf{f}_\star | X, \mathbf{y}, X_\star] &= K(X_\star, X)[K(X, X) + \sigma^2 I]^{-1} \mathbf{y} \\ \text{cov}(\mathbf{f}_\star) &= K(X_\star, X_\star) - K(X_\star, X)[K(X, X) + \sigma^2 I]^{-1} K(X, X_\star) \end{aligned}$$

170 Figure 1 shows an example of samples from a squared exponential Gaus-  
 171 sian Process prior and the updated mean functions after some points have  
 172 been observed.

173 If we look at how to generate predictions for single points, we get the  
 174 following equation of the expectation for new points.

$$\mathbf{f}_\star(\mathbf{x}) = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}_\star)$$

175 with  $\alpha = (K(X, X) + \sigma^2 I)^{-1} \mathbf{y}$ . What this equation tells us is that a  
 176 Gaussian Process can be written as a sum of basis functions. This means  
 177 that a potentially infinitely parametrized model boils down to a finite sum  
 178 when making predictions. This sum only depends on the chosen kernel and  
 179 the data observed thus far (Kac and Siebert, 1947).

180 This is also why Gaussian Process regression is referred to as non-  
 181 parametric. It is not the case that this regression approach has no param-  
 182 eters. Instead, it has potentially infinitely many parameters (parametrized  
 183 by the chosen kernel), but only manifests itself by a finite sum when making




Figure 1: Example of samples from a Gaussian Process prior and posterior. Grey lines indicate samples from the GP. Black dots mark empirical observations. The dark grey line marks the current mean of the GP.

184 predictions. Therefore, Gaussian Process regression is a powerful tool to  
 185 capture many stationary functions. This in turn can be easily applied to  
 186 contexts where the task is to explore or exploit these functions sequentially.

### 187 3.4. General setup

188 Having found a good model to learn functions, that is our first ingredient  
 189 for successful sequential function learning, we now have to find a way to  
 190 smartly explore or exploit the function we are learning over time. Within  
 191 the Gaussian Process approach both pure *exploration* and *exploitation* can  
 192 be viewed as two sides of the same coin. They both take a Gaussian Process

193 to model an underlying unknown function <sup>1</sup> and then estimate the utility  
 194 of all available queries (candidate points from which to sample next) by  
 195 using what is called an acquisition function. An acquisition function can be  
 196 seen as a measurement of usefulness (or utility) that candidate points under  
 197 consideration promise to produce. The approach then goes on to choose as  
 198 the next point the one that currently produces the highest utility. The way  
 199 the general set up works is shown in Algorithm 1.

---

**Algorithm 1** General  $\mathcal{GP}$  optimization Algorithm

---

**Require:** Input space  $\mathcal{D} = \{X, \mathbf{y}\}$ ;  $\mathcal{GP}$ -prior  $\mu_0 = 0, \sigma_0, k$

**for**  $t = 1, 2, \dots$  **do**  
     **Choose**  $x_t^* = \operatorname{argmax} f_{\text{Acquisition}}(\mathbf{x})$   
     **Sample**  $y_t = f(x_t^*) + \epsilon_t$   
     **Perform** Bayesian update for  $\mu_t$  and  $\sigma_t$   
**end for**

---

200     This algorithm starts out with a Gaussian Process distribution over func-  
 201 tions, then assesses the usefulness of the available samples by utilizing the  
 202 acquisition function and selects the point that currently maximizes this func-  
 203 tion. Afterwards, the new output at the chosen sample point is observed,  
 204 the Gaussian Process is updated, and the process starts anew.

---

<sup>1</sup>Sometimes the Gaussian Process here is referred to as a surrogate model (Gramacy and Lee, 2008).

## 205 4. Exploration and Optimal Design

### 206 4.1. Acquisition function

207 The goal in an optimal design setting is to learn an unknown func-  
 208 tion as well and quickly as possible. As said before, this is the same as  
 209 the attempt to maximize information gain over all trials. For a Gaussian,  
 210 the entropy  $H(\mathcal{N}(\mu, \Sigma) = \frac{1}{2} \log |2\pi e \Sigma|$ , which means that in our setting  
 211  $I(\mathbf{y}; f) = \frac{1}{2} \log |I + \sigma^{-2} K|$ , where  $K = [k(x, x')]$ . Even though finding the  
 212 overall information gain maximizer is NP-hard, it can be approximated by an  
 213 efficient greedy algorithm based on Gaussian Processes. If  $F(A) = I(\mathbf{y}_A; f)$ ,  
 214 then this algorithm picks  $x_t = \operatorname{argmax} F(A_{t-1} \cup \{\mathbf{x}\})$ . This in turn can be  
 215 shown to be equivalent to the following acquisition function

$$f_{\text{Acquisition}}(\mathbf{x}_t) = \operatorname{argmax} \sigma_{t-1}(\mathbf{x}) \quad (2)$$


What this algorithm does, is to start out with a Gaussian Process prior,  
 and to sequentially sample from the points that currently show the highest  
 uncertainty in a greedy fashion. Even though this algorithm sounds naïve  
 at first, it can actually be shown to converge to at least a constant fraction  
 of the maximum information gainer (Krause, Singh, and Guestrin, 2008).

$$F(A_T) \geq \left(1 - \frac{1}{e}\right) \max F(A) \quad (3)$$

216 This is based on a property of the acquisition function called *submodularity*  
 217 (Krause and Golovin, 2012). Intuitively, submodularity here means that  
 218 information never hurts (it is always helpful to observe more points), but  
 219 also that the usefulness of newly acquired points decreases the more points  
 220 have been sampled before. This diminishing returns property is crucial to

221 show that the greedy algorithm can be successful over all (see Appendix).  
 222 A simple example of the Gaussian Process uncertainty reduction sampler is  
 223 shown in Figure 2 below. We have sampled a function from a Gaussian Pro-  
 224 cess prior (a squared exponential) and let the algorithm select observations  
 225 by picking as the next observation the one that currently has the highest  
 226 standard deviation attached.

Figure 2: GP-uncertainty reduction example. The dark grey line marks the current mean of the GP. The dashed line shows the upper part of the standard deviation. The light grey lines are samples from the GP.



#### 227 4.2. Example: Learning unknown functions

228 In order to demonstrate how Gaussian Process based exploration works,  
 229 let us simulate how the algorithm learns unknown functions and compare it  
 230 to other algorithms. If we assume that we have to learn an unknown function  
 231 as quickly as possible and that the function  $f$  only takes a one-dimensional  
 232 input  $x = [0, 0.1, 0.2, \dots, 10]$  and to which it maps an output  $y$ , then we can

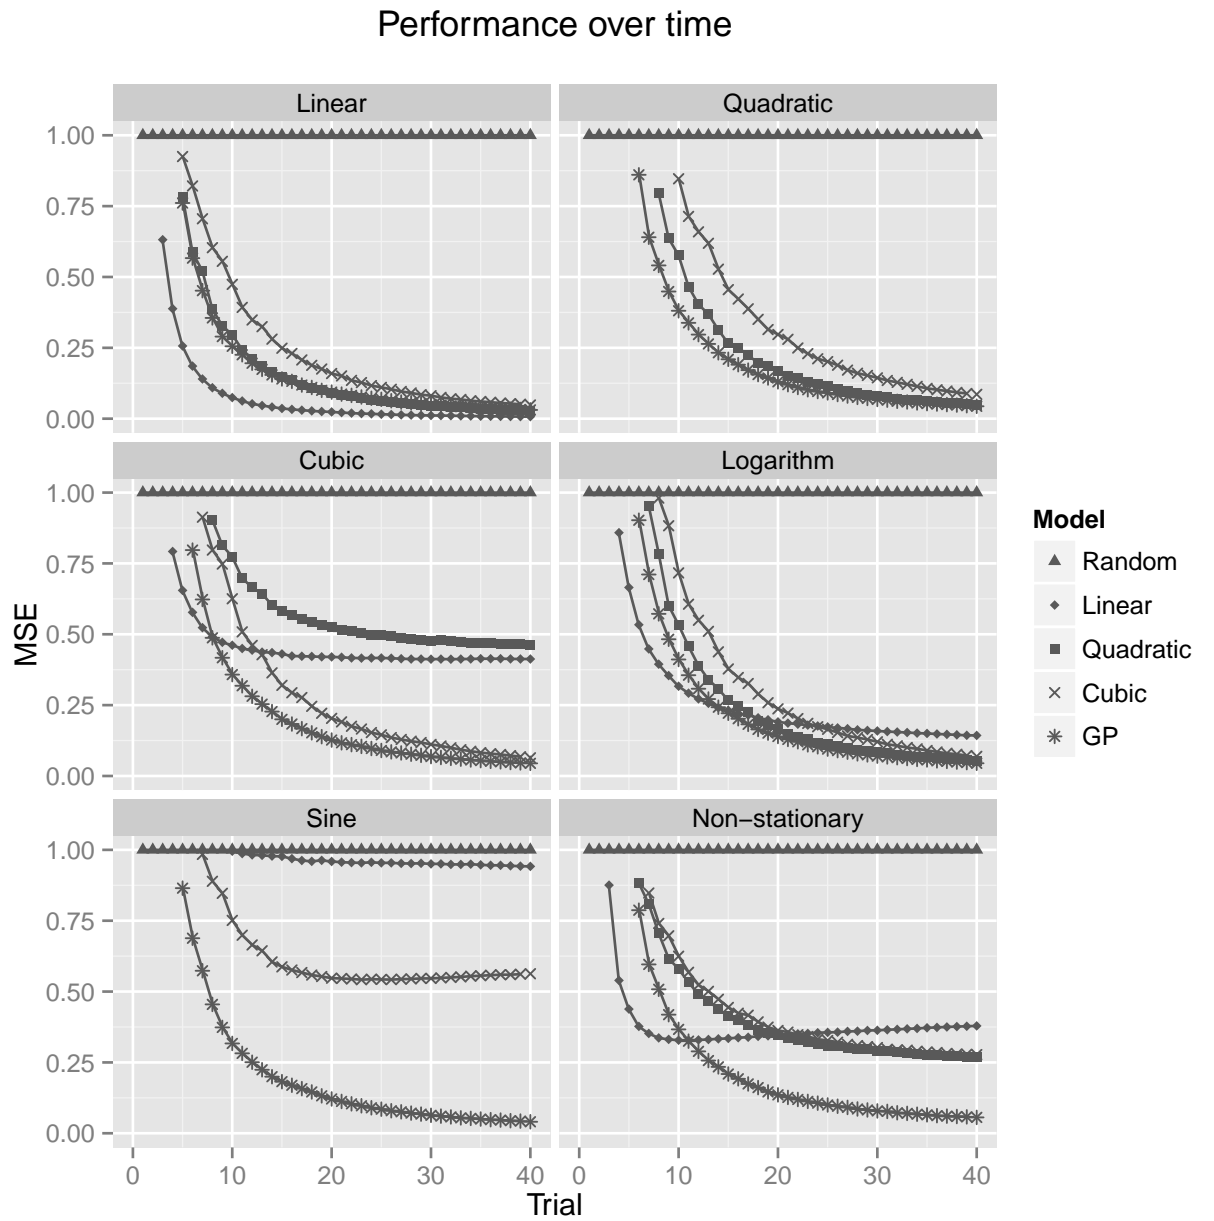
233 set up different functions that the models can learn actively over time. As  
234 the GP is considered to learn many different functions well, we will test the  
235 following different functions: a linear, quadratic, cubic, logarithmic, sine,  
236 and a non-stationary function (see Appendix for details).

237 We have deliberately chosen different parametric forms to assess the  
238 learning methods. The models we have used to learn the different function  
239 are a linear, a quadratic, a cubic, and a Gaussian Process (with a squared  
240 exponential kernel) regression. Each model was set up to learn the under-  
241 lying function by picking as the next observation the one that currently has  
242 the highest uncertainty (standard deviation of the predicted mean). We let  
243 each model run 100 times for each underlying function and averaged the  
244 mean squared error over the whole discretized input space for each step.  
245 Results are shown in Figure 3.

246 It can be seen that the Gaussian Process model learns all functions both  
247 efficiently and well. Only in the cases in which the used learning function  
248 is indeed the same as the learning function (for example, using a linear  
249 function to learn an underlying linear function), does another model learn  
250 faster than the Gaussian Process.

251 This means that Gaussian Processes are especially useful in cases where  
252 the underlying function is not known or can be ignored. For example, one  
253 could easily use Gaussian Processes to learn participants' utility function  
254 over different experiments or simply use them to generate stimuli that are  
255 most informative overall.

Figure 3: GP-uncertainty reduction example.





## 256 5. Exploration-Exploitation and Bayesian Optimization

In an exploration-exploitation scenario the goal is to find the argument to a function that produces the maximum output as quickly as possible. One way to measure the quality of this search process is to quantify regret. Regret is the distance between the output of the currently chosen argument and the maximum output.

$$r(\mathbf{x}) = f(\mathbf{x}^*) - f(\mathbf{x}) \quad (4)$$

The goal then is to minimize regret over all available trials.

$$\min \sum_t^T r(x_t) = \max \sum_t^T f(x_t) \quad (5)$$


Again, finding the global function maximizer is NP hard. That is finding the minimum sequence of queries that leads to the lowest regret overall is almost impossible. However, there is again a greedy trick one can apply in this scenario. This trick is based on redefining the function maximization problem as a bandit task. In a bandit tasks the goal is to maximize output by playing the right arm out of many available arms (it is named after the one armed-bandits that can be found in casinos). As the tasks is to maximize output of a function, the discretized input points to that functions are seen as arms, that are correlated in dependency of the underlying covariance kernel. Taking this perspective, one easy way to approach these kind of problems is the following acquisition function called Upper Confidence Band sampling.

$$f_{\text{Acquisition}}(\mathbf{x}) = \mu_{t-1}(\mathbf{x}) + k\sigma_{t-1}(\mathbf{x}) \quad (6)$$

257 Upper Confidence Band sampling (UCB) plays the arm that currently shows  
 258 the highest upper confidence interval. This strategy is sometimes called

naive optimism in the face of uncertainty. Intuitively, the upper confidence band is determined by two factors, the current estimate of the mean at a particular point (the higher the estimate, the higher the band) and the uncertainty attached to that estimate (the higher the uncertainty, the higher the band). Therefore, the UCB algorithm trades off naturally between expectations and uncertainties. An example of how the UCB works is shown in Figure 4. Again, we have sampled a function from a Gaussian Process prior and have applied the algorithm to this function.

Figure 4: GP-UCB example. The dark grey line marks the current mean of the GP. The dashed line marks the GP’s upper confidence bound. The light grey lines are samples from the GP.



Even though the greedy UCB strategy is again naïve, it can be shown that its regret is sublinear, again based on the submodularity of the overall information gain (Srinivas, Krause, Kakade, and Seeger, 2009).

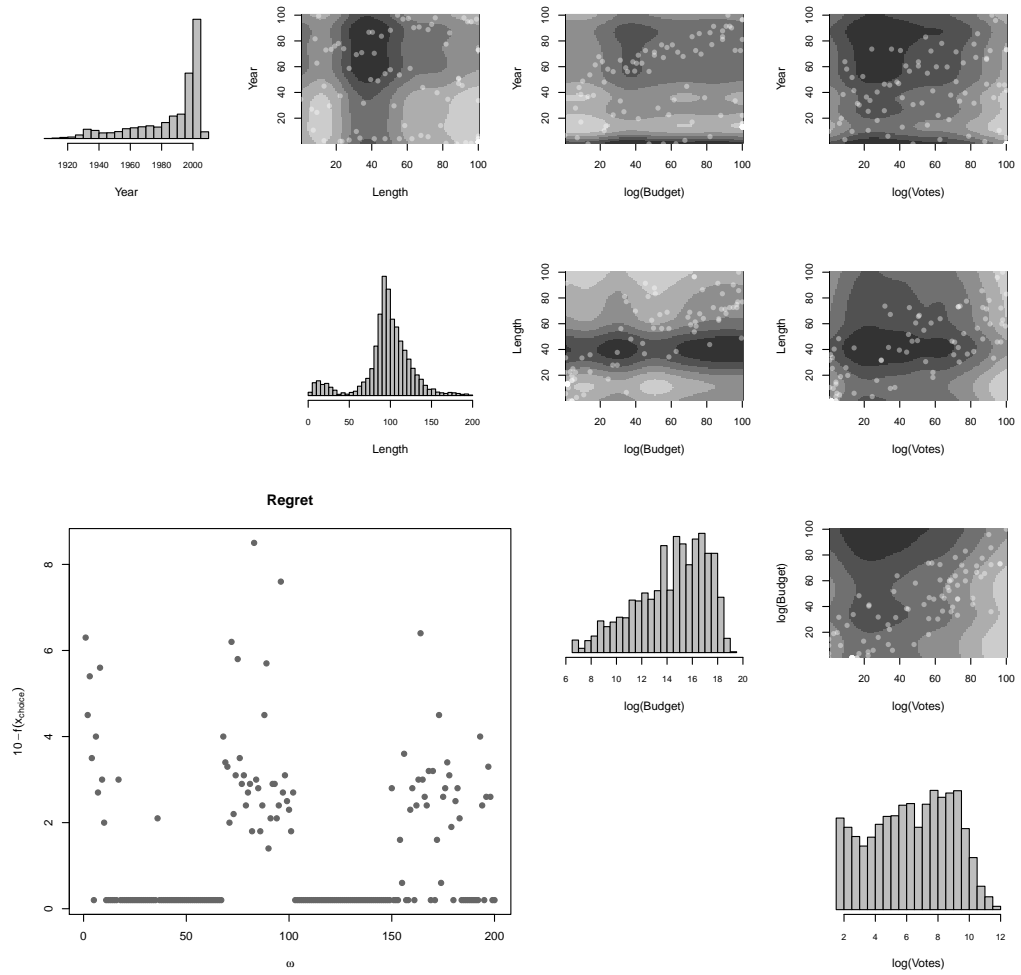
270 *5.1. GP-UCB Example: Choosing movies*

271 We use the imdb-movie data as an example for doing GP-UCB exploratio-  
272 exploitation. For this, we extracted 5141 movies from the data base, includ-  
273 ing the year they appeared, the budget that was used to make them, their  
274 length, as well as how many people had evaluated the movie on the imdb.  
275 The dependent variable was the actual imdb score and we set up a GP-UCB  
276 with a squared exponential kernel, set  $\beta = 3$  and let the algorithm pick 200  
277 movies sequentially. Results are shown in Figure 5.

278 It can be seen that the algorithm quickly starts picking movies that  
279 produce low to zero regret. However, it still keeps exploring other movies  
280 from time to time, just to always come back to recommending movies with  
281 really high scores. Additionally, it explores all of the given input areas for  
282 quite some time, just to find good areas in the end.

Figure 5: GP-UCB example.

### Gaussian Process Optimization



## 283 6. Safe Exploration-Exploitation

284 Sometimes an experimentalist’s goal might not only be to learn or max-  
285 imize a function, but additional requirements can also be important. One  
286 such requirement, for example, can be to avoid certain outputs as much as  
287 possible. One example for such a scenario is excitatory stimulation, espe-  
288 cially in a physiologically setting. Imagine a medical scenario, where the  
289 task is to stimulate the spinal chord in such a way that certain movements  
290 are achieved (Desautels, 2014). Here, it is important to stimulate the spinal  
291 chord such that optimal recovery is achieved, but not too much as this might  
292 lead to painful reactions within the patients. Again, Gaussian Process opti-  
293 mization methods can be used here to learn the underlying function of what  
294 stimulations lead to what strength of reaction. However, an additional re-  
295 quirement now is to avoid these expectedly painful areas. It turns out that  
296 there is a smart way to adapt the GP-UCB approach described above while  
297 accommodating for this additional requirement. This is achieved by an al-  
298 gorithm called SafeOpt (Sui, Gotovos, Burdick, and Krause, 2015). Leaving  
299 the detailed technical explanation of this algorithm for the interested the  
300 reader to look up, this algorithm basically works by trading-off two different  
301 things. Firstly, it keeps a set of safe options it considers to be above the  
302 given threshold and tries to expand this set as much as it can. Secondly,  
303 it maintains a set of potential maximizers that, if used as an input, would  
304 potentially achieve the highest output. It then choses as the next point, a  
305 point within the intersection of these two sets that has the highest predictive  
306 variance.


307 *6.1. Example: Stimulus Optimisation*

308     Imagine that you have to optimize a stimulus that can be described by  
309     two dimensions  $X_1$  and  $X_2$ . This means you have to find the maximum  
310     output  $y$ . However, you also want to avoid sampling values below 0 at all  
311     costs. For this we have drawn a two dimensional function from a Gaussian  
312     Process prior with a squared exponential kernel. This can be seen as a  
313     similar to the case where one want to present stimuli to participants, but  
314     make sure that participants never react with an intensity below a certain  
315     threshold. Results are shown in Figure 6.

316     It can be seen that the SafeOpt algorithm finds the maximum of the  
317     function well, but also tries to expand the space of possible inputs more  
318     and more. At the same time, the algorithm does not sample from the white  
319     area (values below 0) for a single time. We therefore think that this algo-  
320     rithm could potentially applied in many optimal design settings that require  
321     additional constraints.

Figure 6: GP-SafeOpt example. White areas represent areas below 0. The black crosses show where the SafeOpt algorithm has sampled.

### GP Safe Optimization



## 322 7. Gaussian Processes and cognition

323 As we have seen, Gaussian Processes (in combination with smart ac-  
324 quisition functions) are a powerful tool to learn and maximize unknown  
325 functions. However, they might also be applied in a different, even more  
326 psychological context, that is as as a model for human cognition within  
327 the tradition of Bayesian cognitive science (Chater and Oaksford, 2008).  
328 Non-parametric Bayesian approaches have been used before to describe as-  
329 sociative learning (Gershman and Niv, 2012) and categorization learning  
330 (Kemp, Tenenbaum, Griffiths, Yamada, and Ueda, 2006), among others.  
331 Recently, Lucas, Griffiths, Williams, and Kalish (2015) have proposed to  
332 use Gaussian Processes as a rational model of human function learning for  
333 passive learning tasks. Schulz, Tenenbaum, Reshef, Speekenbrink, and Ger-  
334 shman (2015c) used Gaussian Processes to assess participants judgments  
335 about the predictability of functions in dependency of the smoothness of  
336 the underlying kernel. Taking recent developments within the active learn-  
337 ing community into account, we expect Gaussian Processes to soon become  
338 a fruitful descriptive model for human cognition in many different domains  
339 where participants have to act actively, that is to select information sequen-  
340 tially. In a first attempt, we have found that participants' behavior can be  
341 well described by Gaussian Process algorithms when the task is to learn and  
342 maximize a function both in a static (Schulz, Konstantinidis, and Speeken-  
343 brink, 2015b) as well as in a dynamic environment (Schulz, Konstantinidis,  
344 and & Speekenbrink, 2015a). However, given GP's expressiveness and the  
345 mathematical guarantees that come with them, we expect them to be used  
346 more frequently as a model for human cognition in the near future and hope  
347 that this tutorial can help people to apply them more often.



348

349       Here, it is not only the assumption of not having to choose a parametric  
350 shape a priori, but letting the data speak directly, that makes these models  
351 powerful as a psychological model. It is also the attached measure of un-  
352 certainty that comes for free when doing computations with Gaussian Pro-  
353 cesses, a characteristic more and more used in numerics (Hennig, Osborne,  
354 and Girolami, 2015) and optimization problems (Hennig and Kiefel, 2013).  
355 The resulting uncertainty of computation can then also be propagated be-  
356 tween different systems easily (Damianou and Lawrence, 2012). Using this  
357 fact, one could build of a model that involves control, learning, and opti-  
358 mization and in which uncertainty is modeled over the whole system, sort  
359 of a Bayesian Cognitivism approach.

## 360 **8. Discussion**

361       This tutorial has introduced Gaussian Process regression as a general  
362 purpose inference engine to learn about (explore) and maximize (exploit)  
363 unknown functions. Gaussian Processes have been introduced mathemati-  
364 cally and the greedy variance minimization exploration algorithm as well as  
365 the upper confidence bound sampling method for exploration-exploitation  
366 scenarios have been introduced. Within a simulated exploration experiment  
367 we have shown how Gaussian Processes can be used to efficiently learn about  
368 unknown functions in an optimal design set-up and gradually presented in-  
369 formative stimuli. Within a parameter tuning example, we have shown how  
370 GP-UCB can outperform other commonly used methods when the goal is  
371 to optimize hyper-parameters of a neural network. Additionally, we have  
372 talked about introducing additional requirements to the used acquisition

373 function and have shown one example within a scenario, where the task was  
374 to show the optimal excitatory stimulus while avoiding some areas of the in-  
375 put space. Finally, we have talked about utilizing Gaussian Process models  
376 as actual models of cognition in the hope that they will be picked up and  
377 used more frequently in the near future.

378

379     Of course a tutorial like this can never be fully comprehensive. Therefore,  
380 let us briefly point out some things that we have not covered here. First of  
381 all, using the Gaussian Process regression approach as we have parametrized  
382 here is only one possible approach towards regression problems. In fact, it  
383 can be shown that many standard Bayesian regression approaches can be  
384 re-parametrized to be equivalent to Gaussian Process regression, given spe-  
385 cific assumptions about the kernel (Duvenaud, Lloyd, Grosse, Tenenbaum,  
386 and Ghahramani, 2013). The two chosen utility functions here also are just  
387 two options out of a pool of different acquisition functions. Another com-  
388 monly used acquisition function for the pure exploration case is the attempt  
389 to minimize the expected variance summed up over the whole input space  
390 (Gramacy and Apley, 2014). Even though this method tends to sample less  
391 from the boundary cases as compared to the algorithm introduced here, it  
392 can be hard to compute, especially if the input space is large. There exist  
393 many different acquisition functions in the exploration-exploitation context,  
394 that are mostly discussed under the umbrella term Bayesian optimization  
395 (de Freitas, Smola, and Zoghi, 2012). One other common acquisition func-  
396 tion that is frequently used here is the expected probability of improvement  
397 (Mockus, 2012), which choses as the next point the one that promises to  
398 have the highest likelihood of improving the currently expected maximum.  
399 Again, this method of assessing candidate points can be computationally ex-

400 pensive and mathematical guarantees are only given for the UCB algorithm.  
 401 Last not least, there are also many different acquisition functions with ad-  
 402 ditional constraints available. Important to mention hereby are Bayesian  
 403 black box optimization with additional constraints, the estimation of level  
 404 sets, as well as Bayesian Quadrature-based algorithms.

405

406 Here we have introduced Gaussian Process-based algorithms to explore  
 407 and exploit unknown functions. We hope that this tutorial will help others  
 408 to pick up these methods and that their usage will gradually become more  
 409 common within psychology.

## 410 9. Software packages

411 Table 1 below contains further Gaussian Process software pointers.

Table 1: Gaussian Process Software Packages

Name	Algorithm	Language	Author
GPML	GP Toolbox	Matlab	Rasmussen & Williams
SFO	Submodular Optimization	Matlab	Krause
GPy	GP Toolbox	Python	Sheffield ML Group
gptk	GP Toolbox	R	Lawrence
tgp	Tree GPs, GP regression	R	Gramacy & Taddy

## 412 References

- 413 Borji, A., Itti, L., 2013. Bayesian optimization explains human active search.  
 414 In: Advances in neural information processing systems. pp. 55–63.
- 415 Brochu, E., Cora, V. M., De Freitas, N., 2010. A tutorial on bayesian  
 416 optimization of expensive cost functions, with application to active

417 user modeling and hierarchical reinforcement learning. arXiv preprint  
418 arXiv:1012.2599.

419 Cavagnaro, D. R., Aranovich, G. J., McClure, S. M., Pitt, M. A., Myung,  
420 J. I., 2014. On the functional form of temporal discounting: An optimized  
421 adaptive test.

422 Cavagnaro, D. R., Myung, J. I., Pitt, M. A., Kujala, J. V., 2010. Adap-  
423 tive design optimization: A mutual information-based approach to model  
424 discrimination in cognitive science. *Neural computation* 22 (4), 887–905.

425 Chater, N., Oaksford, M., 2008. The probabilistic mind: Prospects for  
426 Bayesian cognitive science. Oxford University Press.

427 Damianou, A. C., Lawrence, N. D., 2012. Deep gaussian processes. arXiv  
428 preprint arXiv:1211.0358.

429 Daunizeau, J., Preuschoff, K., Friston, K., Stephan, K., 2011. Optimizing  
430 experimental design for comparing models of brain function. *PLoS Com-  
431 put Biol* 7 (11), e1002280.

432 de Freitas, N., Smola, A., Zoghi, M., 2012. Regret bounds for deterministic  
433 gaussian process bandits. arXiv preprint arXiv:1203.2177.

434 Desautels, T. A., 2014. Spinal cord injury therapy through active learning.  
435 Ph.D. thesis, California Institute of Technology.

436 Duvenaud, D., Lloyd, J. R., Grosse, R., Tenenbaum, J. B., Ghahramani, Z.,  
437 2013. Structure discovery in nonparametric regression through composi-  
438 tional kernel search. arXiv preprint arXiv:1302.4922.

439 Gershman, S. J., Blei, D. M., 2012. A tutorial on bayesian nonparametric  
440 models. *Journal of Mathematical Psychology* 56 (1), 1–12.

441 Gershman, S. J., Niv, Y., 2012. Exploring a latent cause theory of classical  
442 conditioning. *Learning & behavior* 40 (3), 255–268.

443 Goos, P., Jones, B., 2011. Optimal design of experiments: a case study  
444 approach. John Wiley & Sons.

445 Gramacy, R. B., Apley, D. W., 2014. Local gaussian process approximation  
446 for large computer experiments. *Journal of Computational and Graphical*  
447 *Statistics* (just-accepted), 1–28.

448 Gramacy, R. B., Lee, H. K., 2008. Bayesian treed gaussian process mod-  
449 els with an application to computer modeling. *Journal of the American*  
450 *Statistical Association* 103 (483).

451 Hennig, P., Kiefel, M., 2013. Quasi-newton methods: A new direction. *The*  
452 *Journal of Machine Learning Research* 14 (1), 843–865.

453 Hennig, P., Osborne, M. A., Girolami, M., 2015. Probabilistic numerics and  
454 uncertainty in computations. *Proc. R. Soc. A* 471 (2179), 20150142.

455 Jäkel, F., Schölkopf, B., Wichmann, F. A., 2007. A tutorial on kernel meth-  
456 ods for categorization. *Journal of Mathematical Psychology* 51 (6), 343–  
457 358.

458 Kac, M., Siebert, A., 1947. An explicit representation of a stationary gaus-  
459 sian process. *The Annals of Mathematical Statistics*, 438–442.

460 Kemp, C., Tenenbaum, J. B., Griffiths, T. L., Yamada, T., Ueda, N., 2006.

- 461     Learning systems of concepts with an infinite relational model. In: AAAI.  
462     Vol. 3. p. 5.
- 463     Ko, C.-W., Lee, J., Queyranne, M., 1995. An exact algorithm for maximum  
464     entropy sampling. *Operations Research* 43 (4), 684–691.
- 465     Krause, A., Golovin, D., 2012. Submodular function maximization.  
466     *Tractability: Practical Approaches to Hard Problems* 3, 19.
- 467     Krause, A., Singh, A., Guestrin, C., 2008. Near-optimal sensor placements  
468     in gaussian processes: Theory, efficient algorithms and empirical studies.  
469     *The Journal of Machine Learning Research* 9, 235–284.
- 470     Lucas, C. G., Griffiths, T. L., Williams, J. J., Kalish, M. L., 2015. A rational  
471     model of function learning. *Psychonomic bulletin & review*, 1–23.
- 472     Meder, B., Nelson, J. D., 2012. Information search with situation-specific  
473     reward functions. *Judgment and Decision Making* 7 (2), 119–148.
- 474     Mockus, J., 2010. Bayesian Heuristic approach to discrete and global opti-  
475     mization: Algorithms, visualization, software, and applications. Springer-  
476     Verlag.
- 477     Mockus, J., 2012. Bayesian approach to global optimization: theory and  
478     applications. Vol. 37. Springer Science & Business Media.
- 479     Myung, J. I., Pitt, M. A., 2009. Optimal experimental design for model  
480     discrimination. *Psychological review* 116 (3), 499.
- 481     Rasmussen, C. E., 2006. Gaussian processes for machine learning.

482 Schulz, E., Konstantinidis, E., & Speekenbrink, M., 2015a. Learning and  
 483 decisions in contextual multi-armed bandit tasks. In: Proceedings of the  
 484 Thirty-Seventh Annual Conference of the Cognitive Science Society.

485 Schulz, E., Konstantinidis, E., Speekenbrink, M., 2015b. Exploration-  
 486 exploitation in a contextual multi-armed bandit task. In: International  
 487 Conference on Cognitive Modeling. pp. 118–123.

488 Schulz, E., Tenenbaum, J. B., Reshef, D. N., Speekenbrink, M., Gersh-  
 489 man, S. J., 2015c. Assessing the perceived predictability of functions. In:  
 490 Proceedings of the Thirty-Seventh Annual Conference of the Cognitive  
 491 Science Society.

492 Snoek, J., Larochelle, H., Adams, R. P., 2012. Practical bayesian optimiza-  
 493 tion of machine learning algorithms. In: Advances in neural information  
 494 processing systems. pp. 2951–2959.

495 Srinivas, N., Krause, A., Kakade, S. M., Seeger, M., 2009. Gaussian process  
 496 optimization in the bandit setting: No regret and experimental design.  
 497 arXiv preprint arXiv:0912.3995.

498 Sui, Y., Gotovos, A., Burdick, J., Krause, A., 2015. Safe exploration for  
 499 optimization with gaussian processes. In: Proceedings of the 32nd Inter-  
 500 national Conference on Machine Learning (ICML-15). pp. 997–1005.

501 Wetzels, R., Vandekerckhove, J., Tuerlinckx, F., Wagenmakers, E.-J., 2010.  
 502 Bayesian parameter estimation in the expectancy valence model of the  
 503 iowa gambling task. *Journal of Mathematical Psychology* 54 (1), 14–27.

504 Williams, C. K., 1998. Prediction with gaussian processes: From linear re-

505      gression to linear prediction and beyond. In: Learning in graphical models.  
506      Springer, pp. 599–621.